

Recall Features \underline{x} (p-variate) + response y (let's say continuous), neural network model is a model for $f(\underline{x})$ in

$$\hat{y} = f(\underline{x}) + \varepsilon$$

where

$$f(\underline{x}) = \beta_0 + \sum_{k=1}^{K_L} \beta_k A_k^{(L)}$$

$$A_k^{(L)} = g \left(w_{k0}^{(L)} + \sum_{j=1}^{K_{L-1}} w_{kj}^{(L)} A_j^{(L-1)} \right) \quad k=1, \dots, K_L$$

\vdots

$$A_k^{(1)} = g \left(w_{k0}^{(1)} + \sum_{j=1}^p w_{kj}^{(1)} x_j \right) \quad k=1, \dots, K_1$$

where L hidden layers, the k th of which has K_k activations using the activation function g (usually ReLU).

Note If we set $\underline{x} = (1, x_1, \dots, x_p)^T$ to include the intercept, then we can write

$$A_1^{(0)} = 1 \quad (= \text{bias holder for next level})$$

$$A_1^{(1)} = g\left(w_{10}^{(1)} + \sum_{j=1}^p w_{1j}^{(1)} x_j\right) \Rightarrow g\left(\underbrace{W^{(1)} \underline{x}}_{(p+1) \times 1}\right)$$

$$A_2^{(1)} = g\left(w_{20}^{(1)} + \sum_{j=1}^p w_{2j}^{(1)} x_j\right)$$

$$\vdots$$

$$A_{K_1}^{(1)} = g\left(w_{K_1 0}^{(1)} + \sum_{j=1}^p w_{K_1 j}^{(1)} x_j\right)$$

$(1+K_1) \times (p+1)$
 $(p+1) \times 1$ has a 1 in 1st entry
 apply to each element

$$W^{(1)} = \begin{pmatrix} 1 & 0 & \dots & 0 \\ w_{10}^{(1)} & w_{11}^{(1)} & \dots & w_{1P}^{(1)} \\ w_{20}^{(1)} & w_{21}^{(1)} & \dots & w_{2P}^{(1)} \\ \vdots & \vdots & \ddots & \vdots \\ w_{K0}^{(1)} & w_{K1}^{(1)} & \dots & w_{KP}^{(1)} \end{pmatrix}$$

10.3 Estimation

Estimation can be difficult due to the # of parameters in a row.

One approach: minimize

$$\sum_{i=1}^n (y_i - f(x_i))^2$$

Problem: nonconvex, not guaranteed to have unique soln.

$$R(\underline{\theta}) = \frac{1}{2} \sum_{i=1}^n (y_i - f_{\underline{\theta}}(x_i))^2$$

where $\underline{\theta}$ = vector of (β 's + all w 's)

In practice we use a gradient descent algorithm to estimate $\underline{\theta}$:

① Start with a guess for $\underline{\theta}^{(0)}$ + set $t=0$.

[start at rid N_0]

② Iterate until MSE doesn't improve :

③ Find vector \underline{s} reflecting a small change in $\underline{\theta}$ such that $\underline{\theta}^{t+1} = \underline{\theta}^t + \underline{s}$ and $R(\underline{\theta}^{t+1}) < R(\underline{\theta}^t)$

④ set $t \leftarrow t+1$

The \underline{s} vector is just a scaled gradient:

$$\underline{\theta}^{t+1} = \underline{\theta}^t - \rho \nabla R(\underline{\theta}^t)$$

where $\nabla R(\underline{\theta}^t)$ is gradient of R evaluated at current value,

and $\eta > 0$ is the learning rate, and is typically small.

Note The gradients for NNs end up being straightforward.

Example: single layer feed forward network

$$\begin{aligned} R_i(\underline{\theta}) &= \frac{1}{2} (y_i - f_{\underline{\theta}}(x_i))^2 \\ &= \frac{1}{2} \left(y_i - \beta_0 - \sum_{k=1}^K \beta_k g \left(\underbrace{w_{k0} + \sum_{j=1}^P w_{kj} x_{ij}}_{z_{ik}} \right) \right)^2 \\ &= \frac{1}{2} \left(y_i - \beta_0 - \sum_{k=1}^K \beta_k g(z_{ik}) \right)^2 \end{aligned}$$

$$\frac{\partial R_i(\underline{\theta})}{\partial \beta_k} = \frac{\partial R_i(\underline{\theta})}{\partial f_{\underline{\theta}}(x_i)} \frac{\partial f_{\underline{\theta}}(x_i)}{\partial \beta_k}$$

$$= -(y_i - f_{\theta}(x_i)) \cdot g(z_{ik})$$

$$\frac{\partial R_i(\theta)}{\partial w_{kj}} = \frac{\partial R_i(\theta)}{\partial f_{\theta}(x_i)} \cdot \frac{\partial f_{\theta}(x_i)}{\partial g(z_{ik})} \cdot \frac{\partial g(z_{ik})}{\partial z_{ik}} \cdot \frac{\partial z_{ik}}{\partial w_{kj}}$$

$$= -(y_i - f_{\theta}(x_i)) \cdot \beta_k \cdot g'(z_{ik}) \cdot x_{ij}$$

Each deriv gets the residual $(y_i - f_i)$, which is called backpropagation.

Need some additional tricks b/c θ is high dimensional

- For sample of size n , a smaller random subset of data is used in each gradient step, "mini batch", resulting in stochastic gradient descent.

- Regularization is often done using ~~loss~~ or ridge-like penalty.
- Dropout learning uses a subset of activation nodes at each step, allowing us to focus on only a subset of weights.